

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$
where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top x_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i x_i)$
- Return final weight vector

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$
where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top x_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i x_i)$
- Return final weight vector

Remember:

Prediction = $\text{sgn}(\mathbf{w}^\top \mathbf{x})$

There is typically a bias term also ($\mathbf{w}^\top \mathbf{x} + b$), but the bias may be treated as a constant feature and folded into \mathbf{w}

Footnote: For some algorithms it is mathematically easier to represent False as -1, and at other times, as 0. For the Perceptron algorithm, treat -1 as false and +1 as true.

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$

where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

Mistake on positive: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r \mathbf{x}_i$
Mistake on negative: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r \mathbf{x}_i$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top \mathbf{x}_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i \mathbf{x}_i)$
- Return final weight vector

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$

where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

Mistake on positive: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r \mathbf{x}_i$
Mistake on negative: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r \mathbf{x}_i$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top \mathbf{x}_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i \mathbf{x}_i)$
- Return final weight vector

r is the learning rate, a small positive number less than 1

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$

where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top x_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i x_i)$
- Return final weight vector

Mistake on positive: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r x_i$
Mistake on negative: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r x_i$

r is the learning rate, a small positive number less than 1

Update only on error. A mistake-driven algorithm

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$

where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top x_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i x_i)$
- Return final weight vector

Mistake on positive: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r x_i$
Mistake on negative: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r x_i$

r is the learning rate, a small positive number less than 1

Update only on error. A mistake-driven algorithm

This is the simplest version. We will see more robust versions at the end

The Perceptron algorithm

Input: A sequence of training examples $(x_1, y_1), (x_2, y_2), \dots$

where all $x_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

- Initialize $\mathbf{w}_0 = \mathbf{0} \in \mathbb{R}^n$
- For each training example (x_i, y_i) :
 - Predict $y' = \text{sgn}(\mathbf{w}_t^\top x_i)$
 - If $y_i \neq y'$:
 - Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r (y_i x_i)$
- Return final weight vector

Mistake on positive: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r x_i$
Mistake on negative: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r x_i$

r is the learning rate, a small positive number less than 1

Update only on error. A mistake-driven algorithm

This is the simplest version. We will see more robust versions at the end

Mistake can be written as $y_i \mathbf{w}_t^\top x_i \leq 0$

Intuition behind the update

Mistake on positive: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r \mathbf{x}_i$

Mistake on negative: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r \mathbf{x}_i$

Suppose we have made a mistake on a positive example

That is, $y = +1$ and $\mathbf{w}_t^\top \mathbf{x} < 0$

Call the **new weight vector** \mathbf{w}_{t+1} .

The update means $\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{x}$ (say $r = 1$)

The **new dot product** will be $\mathbf{w}_{t+1}^\top \mathbf{x} = (\mathbf{w}_t + \mathbf{x})^\top \mathbf{x} = \mathbf{w}_t^\top \mathbf{x} + \mathbf{x}^\top \mathbf{x} \geq \mathbf{w}_t^\top \mathbf{x}$

For a positive example, the Perceptron update will increase the score assigned to the same input

Similar reasoning for negative examples

1. The “standard” algorithm

Given a training set $D = \{(\mathbf{x}_i, y_i)\}$, $\mathbf{x}_i \in \mathbb{R}^n$, $y_i \in \{-1, 1\}$

1. Initialize $\mathbf{w} = \mathbf{0} \in \mathbb{R}^n$
2. For epoch = 1 ... T:
 1. Shuffle the data
 2. For each training example $(\mathbf{x}_i, y_i) \in D$:
 - If $y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$, update $\mathbf{w} \leftarrow \mathbf{w} + \underline{r} y_i \mathbf{x}_i$
3. Return \mathbf{w}

Prediction: $\text{sgn}(\mathbf{w}^\top \mathbf{x})$